

# Introduction to tidytext

Data Wrangling and Husbandry  
2/25/2019



# Textual Analysis

There are many analyses of text that one might want to do, including

- Frequencies of words
- Sentiment of words
- Discovery of topics

Until recently, the main formats have been

- Raw strings
- Corpus: raw string content with additional metadata and details
- Document-term matrix

Julia Silge and David Robinson have introduced the `tidytext` package and new [book](#) which offers the tidytext format. That format is simply a table with one-term-per-row.

(Installing the `tidytext` package may take a while if you install from source, since it requires installing many additional packages.)

This is not a replacement for full-fledged natural language processing, but is an easy entry point for many analyses.



# An extended example (from Tidy Text Mining with R )

We will make a frequency table of words in *Pride and Prejudice* and follow that with sentiment analysis of all of Jane Austen's novels.

```
library(janeaustenr)
prideprejudice[1:11]

## [1] "PRIDE AND PREJUDICE"
## [2] ""
## [3] "By Jane Austen"
## [4] ""
## [5] ""
## [6] ""
## [7] "Chapter 1"
## [8] ""
## [9] ""
## [10] "It is a truth universally acknowledged, that
## [11] "of a good fortune, must be in want of a wife"
```

```
library(stringr)
prideprejudice.tbl <- tibble(text = prideprejudice) %
mutate(linenumber = row_number(),
       chapter = cumsum(str_detect(text, "(?i)^chap
prideprejudice.tbl[10:15,]
```

```
## # A tibble: 6 x 3
##   text
##   <chr>
## 1 It is a truth universally acknowledged, that a s
## 2 of a good fortune, must be in want of a wife.
## 3 ""
## 4 However little known the feelings or views of su
## 5 first entering a neighbourhood, this truth is so
## 6 of the surrounding families, that he is consider
```





# unnest\_tokens()

We can move from strings to words with the `unnest_tokens()` function (which by default changes everything to lower case).

```
library(tidytext)
tidy_prideprejudice <- prideprejudice.tbl %>%
  unnest_tokens(word, text)
```

```
tidy_prideprejudice[9:14,]
```

```
## # A tibble: 6 x 3
##   linenumber chapter word
##       <int>    <int> <chr>
## 1         10         1 it
## 2         10         1 is
## 3         10         1 a
## 4         10         1 truth
## 5         10         1 universally
## 6         10         1 acknowledged
```

Now that the data is tidy, we can use regular tools from the tidyverse. For example, what are the most common words?

```
tidy_prideprejudice %>%  
  count(word, sort = TRUE)
```

```
## # A tibble: 6,538 x 2  
##   word      n  
##   <chr> <int>  
## 1 the     4331  
## 2 to      4162  
## 3 of      3610  
## 4 and     3585  
## 5 her     2203  
## 6 i       2065  
## 7 a       1954  
## 8 in      1880  
## 9 was     1843  
## 10 she    1695  
## # ... with 6,528 more rows
```

Hmm, that's not quite what we want. We can eliminate what are known as *stop words*.

```
data(stop_words)
sample_n(stop_words, 10)
```

```
## # A tibble: 10 x 2
##   word      lexicon
##   <chr>    <chr>
## 1 needs    SMART
## 2 ever     SMART
## 3 seems    SMART
## 4 ourselves snowball
## 5 can      onix
## 6 serious  SMART
## 7 regardless SMART
## 8 further  SMART
## 9 immediate SMART
## 10 exactly SMART
```

```
tidy_prideprejudice %>%  
  anti_join(stop_words) %>%  
  count(word, sort = TRUE)
```

```
## Joining, by = "word"
```

```
## # A tibble: 6,009 x 2
```

```
##   word      n
```

```
##   <chr>    <int>
```

```
## 1 elizabeth 597
```

```
## 2 darcy     373
```

```
## 3 bennet    294
```

```
## 4 miss      283
```

```
## 5 jane      264
```

```
## 6 bingley   257
```

```
## 7 time      203
```

```
## 8 lady      183
```


```
## 9 sister    180
```

```
## 10 wickham  162
```

```
## # ... with 5,999 more rows
```

```
library(ggplot2)

tidy_prideprejudice %>%
  anti_join(stop_words) %>%
  count(word, sort = TRUE) %>%
  filter(n > 150) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n)) +
  geom_bar(stat = "identity") +
  xlab(NULL) + coord_flip()
```







# N-grams

What about pairs or triples of words? These are known as “n-grams”, and can be obtained by an optional argument to `unnest_tokens()`.

```
tidy_prideprejudice_bigram <- prideprejudice.tbl %>%
  unnest_tokens(bigram, text, token = "ngrams", n = 2)
tidy_prideprejudice_bigram
```

```
## # A tibble: 114,045 x 3
##   linenumber chapter bigram
##       <int>    <int> <chr>
## 1         1         0 pride and
## 2         1         0 and prejudice
## 3         2         0 <NA>
## 4         3         0 by jane
## 5         3         0 jane austen
## 6         4         0 <NA>
## 7         5         0 <NA>
## 8         6         0 <NA>
## 9         7         1 chapter 1
## 10        8         1 <NA>
## # ... with 114,035 more rows
```

We can't just drop stop words, since that will mess up the pairs.

Instead, we can split out the words in a pair and delete the pairs with a stop word in them

```
prideprejudice_bigrams_separated <- tidy_prideprejudice_bigrams %>%  
  separate(bigram, c("word1", "word2"), sep = " ")
```

```
prideprejudice_bigrams_filtered <- prideprejudice_bigrams_separated %>%  
  filter(!word1 %in% stop_words$word) %>%  
  filter(!word2 %in% stop_words$word)
```

## prideprejudice\_bigrams\_filtered

```
## # A tibble: 8,835 x 4
##   linenumber chapter word1      word2
##       <int>   <int> <chr>    <chr>
## 1         2       0 <NA>    <NA>
## 2         3       0 jane     austen
## 3         4       0 <NA>    <NA>
## 4         5       0 <NA>    <NA>
## 5         6       0 <NA>    <NA>
## 6         7       1 chapter  1
## 7         8       1 <NA>    <NA>
## 8         9       1 <NA>    <NA>
## 9        10       1 truth    universally
## 10       10       1 universally acknowledged
## # ... with 8,825 more rows
```

```
prideprejudice_bigrams_filtered %>%  
  count(word1, word2, sort = TRUE)
```

```
## # A tibble: 5,109 x 3  
##   word1    word2      n  
##   <chr>   <chr>   <int>  
## 1 <NA>    <NA>    2556  
## 2 lady    catherine  87  
## 3 miss    bingley    67  
## 4 miss    bennet     52  
## 5 sir     william    35  
## 6 de      bourgh     32  
## 7 miss    darcy      32  
## 8 cried   elizabeth  24  
## 9 colonel forster  23  
## 10 miss   lucas      23  
## # ... with 5,099 more rows
```



# Sentiment Analysis

Sentiment analysis is an attempt to extract the sentiment, or feeling or attitude, from text. The most popular approach is just to match individual words against a lexicon. For later, note that `get_sentiments()` extracts a particular lexicon.

```
sentiments
```

```
## # A tibble: 27,314 x 4
##   word      sentiment lexicon score
##   <chr>      <chr>      <chr>  <int>
## 1 abacus      trust      nrc      NA
## 2 abandon     fear      nrc      NA
## 3 abandon    negative    nrc      NA
## 4 abandon     sadness    nrc      NA
## 5 abandoned   anger      nrc      NA
## 6 abandoned   fear      nrc      NA
## 7 abandoned   negative    nrc      NA
## 8 abandoned     sadness    nrc      NA
## 9 abandonment anger      nrc      NA
## 10 abandonment fear      nrc      NA
## # ... with 27,304 more rows
```

## (All code from Text Mining With R)

```
tidy_books <- austen_books() %>%  
  group_by(book) %>%  
  mutate(linenumber = row_number(),  
         chapter = cumsum(str_detect(text, regex("^cl  
                                     ignore  
  
  ungroup() %>%  
  unnest_tokens(word, text)
```





# It helps to know your subject

```
sentiments %>% filter(word == "miss")
```

```
## # A tibble: 3 x 4
##   word  sentiment lexicon  score
##   <chr> <chr>      <chr>   <int>
## 1 miss  negative  bing      NA
## 2 miss  <NA>      AFINN     -2
## 3 miss  negative  loughran  NA
```

We could carefully drop “miss” when it’s followed by a character name, but we’ll take the sloppier approach and just filter it out.

```
janeaustensentiment <- tidy_books %>%
  filter(word != "miss") %>%
  inner_join(get_sentiments("bing")) %>%
  count(book, index = linenumber %/% 80, sentiment) %>%
  spread(sentiment, n, fill = 0) %>%
  mutate(sentiment = positive - negative)
```

```
## Joining, by = "word"
```


```
janeaustensentiment
```

```
## # A tibble: 920 x 5
```


##	book	index	negative	positive	ser
##	<fct>	<dbl>	<dbl>	<dbl>	
##	1 Sense & Sensibility	0	16	32	
##	2 Sense & Sensibility	1	19	53	
##	3 Sense & Sensibility	2	11	31	
##	4 Sense & Sensibility	3	15	31	
##	5 Sense & Sensibility	4	16	34	
##	6 Sense & Sensibility	5	15	51	
##	7 Sense & Sensibility	6	24	40	
##	8 Sense & Sensibility	7	23	51	
##	9 Sense & Sensibility	8	30	40	
##	10 Sense & Sensibility	9	15	19	
##	# ... with 910 more rows				

```
ggplot(janeaustensentiment, aes(index, sentiment, fill = sentiment)) +  
  geom_bar(alpha = 0.8, stat = "identity", show.legend = FALSE) +  
  facet_wrap(~book, ncol = 3, scales = "free_x")
```

```
ggplot(janeaustensentiment, aes(index, sentiment, fill = book)) +
  geom_bar(alpha = 0.8, stat = "identity") +
  facet_wrap(~book, ncol = 3, scales = "free_x") +
  theme(legend.position = "none")
```



```
ggplot(janeaustensentiment, aes(index, sentiment, fill = book)) +
  geom_bar(alpha = 0.8, stat = "identity", show.legend = FALSE) +
  facet_wrap(~book, ncol = 3, scales = "free_x") +
  theme(legend.position = "none") +
  geom_smooth(span = .15)
```





# In class exercises

1. Use the `gutenbergr` package to download the texts of Wuthering Heights and Jane Eyre.

```
library(gutenbergr)
bronte_books <- gutenberg_download(gutenberg_id = c('
  meta_fields = "title")
```

1. Tidy the downloaded data so that each row is one word.
2. Find the most popular non-stop-words in Jane Eyre
3. Plot the sentiment for the two books (just re-use the code above)