

Intro to Project Management

Data Wrangling and Husbandry

03/09/2020


Project Management Issues

As a project becomes more complex, there are a number of issues that need to be addressed to effectively work on the project

- ▶ The number of files may increase; you need to keep tabs on them all
- ▶ Your data, your code, and your reports may go through multiple versions
- ▶ You may be working with other people who will need to share your data or code
- ▶ You may be returning to the project months later

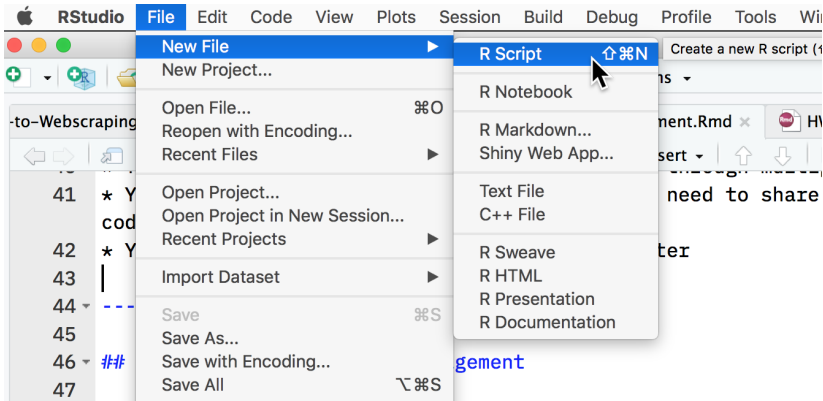
Bare Minimum Project Management

- ▶ Use a separate directory (also known as a folder) for each project
 - ▶ *Everything* needed goes in the directory
- ▶ Use RStudio's "New Project" menu item to either create that directory or to associate an existing directory with a project



Really Minimal Project Management

- ▶ Use a separate directory (also known as a folder) for each project
- ▶ Rstudio project
- ▶ Put the date on every file (including every revision) as a suffix
 - ▶ e.g., mouse-study-2018-03-01.R
- ▶ Make sure there are no R object that you generated in the console window—everything should come from an R script or an RMarkdown file.
 - ▶ I *strongly* recommend running code from an R Script or Notebook file rather than the console.



Minimal Project Management

- ▶ Use a separate directory (also known as a folder) for each project
- ▶ Rstudio project
- ▶ Put the date on every file (including every revision) as a suffix
- ▶ Have separate directories for:
 - ▶ data, or sometimes raw_data and clean_data
 - ▶ code
 - ▶ figures
 - ▶ results
- ▶ (knitr is a little particular about file locations, so I let it write figures and output to where it wants to—I use the figures folder for figures that I explicitly create and save, and similarly for result)
- ▶ Consider including a README.txt or README.md file to describe what the files are
- ▶ Periodically zip the whole thing into a numbered and dated file?

More Sophisticated Project Management

- ▶ Use a Makefile (more later)
- ▶ Use git and github (more later)

(Minimal) Project Management and Reproducibility

- ▶ Reproducible: another group (including future you) should be able to re-create all of the results from just the code and data
- ▶ Requires planning and choice of tools
- ▶ Most essential aspects:
 - ▶ Don't touch the original data—think of it as read-only
 - ▶ Do all of your work via R scripts, notebooks, and RMarkdown files—think of your output as disposable, in the sense that it can be fully recreated from your scripts

My minimal project approach

- ▶ Separate directories for data, code, figures, and results
- ▶ Series of scripts for the work:
 - ▶ file of functions
 - ▶ one or two files to read, clean, and tidy data
 - ▶ export data at this point so that I'll have a tidy version to share in the future
 - ▶ multiple files to do the analysis work, broken up into independent analyses
 - ▶ filenames start with the order number and finish with the date if I'm not using git (e.g., 05-regressions-2016-10-31.R)
- ▶ My final report is usually done via a RMarkdown or Rnw (for mixing R and lots of Latex) file that uses `source()` to run the script files

Workflow versus Product

Jenny Bryan distinguishes between workflow and product:

- ▶ Part of workflow:
 - ▶ The editor you use to write your R code.
 - ▶ The name of your home directory.
 - ▶ The R code you ran before lunch.
- ▶ Part of product:
 - ▶ The raw data.
 - ▶ The R code someone needs to run on your raw data to get your results, including the explicit `library()` calls to load necessary packages.
- ▶ Your project management approach is a way both to
 - ▶ Keep yourself organized
 - ▶ Share your work with others

More on Reproducibility through scripts

- ▶ Don't change the original data.
 - ▶ *No hand-editing*
 - ▶ Make changes to data via script, ideally including comments as to why
 - ▶ If subjects or variables are to be dropped, do so by name, not by position (e.g. `filter(ID != "A12345")`) # withdrew consent, not `example[-258,]`
 - ▶ Ideally put the provenance of the files in the README or in the scripts that load those files

- ▶ Every aspect of cleaning and analysis goes in the scripts. No “repeat this analysis but now using `y2` as the outcome”
 - ▶ If I’ve been working in the console, or making a lot of changes to scripts, I will typically delete everything in the global environment
 - ▶ `rm(list = ls())`, or
 - ▶ RStudio’s “Session > Clear Workspace ...”) to make sure I’m not using hidden definitions
- ▶ Save your seed (e.g., `set.seed(41540130)`) if your analysis involves random number generation.
- ▶ Use relative paths like “../data/example.csv”, not “~/DropBox/OSC/Smith/data/example.csv”
 - ▶ There is a new package [here](#) that makes finding files in your project directory really easy: `here("example.csv")`

- ▶ There are even packages like `packrat` and `checkpoint` to save the versions of the various packages that you used. Recording the results of `sessionInfo()` is sufficient for most purposes, however. Microsoft hosts a CRAN time machine that can be useful for using an older version of a package
- ▶ If you do something repeatedly, you almost always want to do it via a function
 - ▶ the code is easier to follow and you are less like to make a change in one place but not another

- ▶ Refer to rows and variable by name, not position, and don't assume that the rows are in any particular order

```
## example from https://www.stat.ubc.ca/~jenny/STAT545A/blc
## left-hand figure; code contains 44 characters
xyplot(gDat[427:568,5]~log(gDat[427:568,6]))
```

```
## right-hand figure; code contains 202 characters
jYear <- 1967
xyplot(lifeExp ~ gdpPercap, gDat,
       subset = year == jYear, main = paste("year =", jYear),
       group = continent, auto.key = TRUE,
       scales = list(x = list(log = 10, equispaced.log = F))
```


Recommended order for an R script file

1. Load packages
2. Set data independent constants
 - ▶ constants that depend on the data can go in the body of the code, but should be based on functions that invoke the data, e.g. `1:nrow(example)`, not `1:290`
3. Define functions (or `source()` a file with functions)
4. Get to work

Even if I have only a single RMarkdown file I'll put the first 3 steps in the setup chunk.