

# SQL & R

Data Wrangling and Husbandry

03/22/2020



# SQL

- ▶ SQL stands for Structured Query Language.
- ▶ SQL is used to communicate with a database
- ▶ Most commercial relational database management systems rely on SQL for their query language
- ▶ Databases can be very large and are optimized for certain data manipulation tools, so it can be advantageous to use database tools for what they do best and R tools for what they do best

## An aside

One philosophy is to think of R as a flexible user interface. Some tools are built in (e.g., linear models), but other tools interface with other components of the computational world: shell commands, SQL, Spark, and others. Depending on their implementation, they can cut down on the cognitive burden of task switching.

# Why work with a database

- ▶ Your data is already in a database
  - ▶ Relying on, say, csv written out from the database will lead to aging problems
- ▶ You have so much data it won't fit in memory
- ▶ (Besides using a database, there are custom R approaches to the latter situation—the CRAN task view on high performance computing has some details in the “Large memory and out-of-memory data” section)

## Concept mapping

Database	R
Table	Data frame
Field, column	Variable, column
Record, row	Observation, row
to select all cols	blank to select all cols
NULL	NA
single quotes	single or double quotes
not case sensitive	case sensitive

- ▶ There are nice introductions to SQL at <https://www.sqlteaching.com> and <http://2016.padjo.org/tutorials/hello-sqlite-and-sqlite-clients/> , but our goal is not to learn SQL. Instead, our goal is to have R generate and execute SQL commands.

## The dplyr package interfaces with SQL!

In the past, there were database specific functions to connect to databases: \* `src_mysql` connects to a mysql database \* `src_postgres` connects to a postgresql database. \* `src_sqlite` connects to a sqlite database

The current approach is to use the function `DBI::dbConnect()`. The DBI package can be used for database work beyond what dplyr can do.



## DBI::dbConnect()

The arguments to `DBI::dbConnect()` vary from database to database, but the first argument is always the database backend.

- ▶ `RSQLite::SQLite()` for `RSQLite`
- ▶ `RMySQL::MySQL()` for `RMySQL`
- ▶ `RPostgreSQL::PostgreSQL()` for `RPostgreSQL`
- ▶ `odbc::odbc()` for `odbc`
- ▶ `bigrquery::bigquery()` for `BigQuery`

For a remote database, this will look more like

```
my_db <- DBI::dbConnect(RMySQL::MySQL(),  
  host = "example.database.com",  
  user = "example",  
  password = rstudioapi::askForPassword("Database password"  
)
```

You'll have to get help from your database administrator.

## An extended example

Rather than try to tie into an existing database, we will make our own local sqlite database. This is not the way it would work in practice, since if your data fits in R you probably don't need an external database. We will use

```
my_db <- DBI::dbConnect(RSQLite::SQLite(), path = "demo.sqlite")
```

```
library(gapminder)
copy_to(my_db, gapminder, temporary = FALSE, indexes = list(
  "country", "continent", "year"))
```

```
gapminder_sqlite <- tbl(my_db, "gapminder")
```

## dplyr commands for databases

Once you've set up the connection to the database with `src_*`, you can almost treat the database as a collection of data frames.

```
## notice the top matter
```

```
select(gapminder_sqlite, country, continent)
```

```
## # Source:   lazy query [?? x 2]
```

```
## # Database: sqlite 3.30.1 []
```

```
##   country      continent
```

```
##   <chr>         <chr>
```

```
## 1 Afghanistan Asia
```

```
## 2 Afghanistan Asia
```

```
## 3 Afghanistan Asia
```

```
## 4 Afghanistan Asia
```

```
## 5 Afghanistan Asia
```

```
## 6 Afghanistan Asia
```

```
## 7 Afghanistan Asia
```

```
## 8 Afghanistan Asia
```

```
## 9 Afghanistan Asia
```

```
gapminder_sqlite %>% filter(country == "Peru")
```

```
## # Source:   lazy query [?? x 6]
```

```
## # Database: sqlite 3.30.1 []
```

##	country	continent	year	lifeExp	pop	gdpPercap
##	<chr>	<chr>	<int>	<dbl>	<int>	<dbl>
##	1 Peru	Americas	1952	43.9	8025700	3759.
##	2 Peru	Americas	1957	46.3	9146100	4245.
##	3 Peru	Americas	1962	49.1	10516500	4957.
##	4 Peru	Americas	1967	51.4	12132200	5788.
##	5 Peru	Americas	1972	55.4	13954700	5938.
##	6 Peru	Americas	1977	58.4	15990099	6281.
##	7 Peru	Americas	1982	61.4	18125129	6435.
##	8 Peru	Americas	1987	64.1	20195924	6361.
##	9 Peru	Americas	1992	66.5	22430449	4446.
##	10 Peru	Americas	1997	68.4	24748122	5838.
##	# ... with more rows					

In fact, the usual basic commands of `select`, `filter`, `arrange`, `mutate`, and `summarize` work exactly as expected. However, they are translated to SQL and then run on the the database rather than within R

```
(example_query <- gapminder_sqlite %>%  
  filter(year == "2007") %>%  
  group_by(Continent) %>%  
  summarize(lifeExp = mean(lifeExp)))
```

```
## Warning: Missing values are always removed in SQL.  
## Use `mean(x, na.rm = TRUE)` to silence this warning  
## This warning is displayed only once per session.
```

```
## # Source:   lazy query [?? x 2]  
## # Database: sqlite 3.30.1 []  
##   continent lifeExp  
##   <chr>      <dbl>  
## 1 Africa    54.8  
## 2 Americas  73.6  
## 3 Asia      70.7
```

```
show_query(example_query)
```

```
## <SQL>  
## SELECT `Continent`, AVG(`lifeExp`) AS `lifeExp`  
## FROM `gapminder`  
## WHERE (`year` = '2007')  
## GROUP BY `Continent`
```

If you know SQL, you can insert your own code

```
tbl(my_db, sql("SELECT country, continent FROM gapminder"))
```

```
## # Source:   SQL [?? x 2]
## # Database: sqlite 3.30.1 []
##   country      continent
##   <chr>         <chr>
## 1 Afghanistan Asia
## 2 Afghanistan Asia
## 3 Afghanistan Asia
## 4 Afghanistan Asia
## 5 Afghanistan Asia
## 6 Afghanistan Asia
## 7 Afghanistan Asia
## 8 Afghanistan Asia
## 9 Afghanistan Asia
## 10 Afghanistan Asia
## # ... with more rows
```



# Laziness

- ▶ dplyr is *lazy* when it comes to SQL commands
  - ▶ Never brings data into R unless explicitly asked
  - ▶ Collects commands and sends it to the database in one step at the last possible moment
- ▶ In the example below, the SQL code is created but not run

```
lazy_ex <- gapminder_sqlite %>% filter(continent == "Americ  
  select(country, year, pop, gdpPercap) %>%  
  mutate(gdp = pop * gdpPercap) %>% arrange(desc(gdp))
```

If you print the object, dplyr generates the QSL and sends it to the database, but only requests the first 10 rows.

```
lazy_ex
```

```
## # Source:      lazy query [?? x 5]
## # Database:    sqlite 3.30.1 []
## # Ordered by:  desc(gdp)
##   country      year      pop  gdpPercap      gdp
##   <chr>         <int>   <int>    <dbl>    <dbl>
## 1 United States 2007 301139947  42952.  1.29e13
## 2 United States 2002 287675526  39097.  1.12e13
## 3 United States 1997 272911760  35767.  9.76e12
## 4 United States 1992 256894189  32004.  8.22e12
## 5 United States 1987 242803533  29884.  7.26e12
## 6 United States 1982 232187835  25010.  5.81e12
## 7 United States 1977 220239000  24073.  5.30e12
## 8 United States 1972 209896000  21806.  4.58e12
## 9 United States 1967 198712000  19530.  3.88e12
## 10 United States 1962 186538000  16173.  3.02e12
## # ... with more rows
```

To get the full results, use `collect()`

```
collect(lazy_ex)
```

```
## # A tibble: 300 x 5
```

##	country	year	pop	gdpPercap	gdp
##	<chr>	<int>	<int>	<dbl>	<dbl>
##	1 United States	2007	301139947	42952.	1.29e13
##	2 United States	2002	287675526	39097.	1.12e13
##	3 United States	1997	272911760	35767.	9.76e12
##	4 United States	1992	256894189	32004.	8.22e12
##	5 United States	1987	242803533	29884.	7.26e12
##	6 United States	1982	232187835	25010.	5.81e12
##	7 United States	1977	220239000	24073.	5.30e12
##	8 United States	1972	209896000	21806.	4.58e12
##	9 United States	1967	198712000	19530.	3.88e12
##	10 United States	1962	186538000	16173.	3.02e12
##	# ... with 290 more rows				

## Grouping

Compare the following

```
gapminder %>% filter(year == 2007) %>%  
  select(country, continent, lifeExp) %>%  
  group_by(continent) %>% filter(lifeExp == min(lifeExp))
```

```
## # A tibble: 5 x 3  
## # Groups:   continent [5]  
##   country      continent lifeExp  
##   <fct>        <fct>      <dbl>  
## 1 Afghanistan Asia         43.8  
## 2 Haiti        Americas     60.9  
## 3 New Zealand Oceania      80.2  
## 4 Swaziland    Africa       39.6  
## 5 Turkey       Europe      71.8
```

```
gapminder_sqlite %>% filter(year == 2007) %>%  
  select(country, continent, lifeExp)  
%>% group_by(continent) %>% filter(lifeExp == min(lifeExp
```

```
## Error: <text>:3:3: unexpected SPECIAL  
## 2:   select(country, continent, lifeExp)  
## 3:   %>%  
##      ^
```

SQLite has very limited “window” functions—basically, you can summarize but not filter or mutate. PostgreSQL does allow these, and is generally more powerful.

For our example, we can get around the problem by bringing the data to R just before the last filtering step

```
group_filter_ex <- gapminder_sqlite %>%  
  filter(year == 2007) %>%  
  select(country, continent, lifeExp) %>%  
  group_by(continent) %>% collect()  
  
group_filter_ex %>% filter(lifeExp == min(lifeExp))
```

```
## # A tibble: 5 x 3  
## # Groups:   continent [5]  
##   country      continent lifeExp  
##   <chr>         <chr>      <dbl>  
## 1 Afghanistan Asia         43.8  
## 2 Haiti         Americas    60.9  
## 3 New Zealand  Oceania     80.2  
## 4 Swaziland    Africa      39.6  
## 5 Turkey        Europe     71.8
```

## Joins

Joins generally require the data to be in the same location.

You can use `*_join(x, y, copy = TRUE)` to copy `y` into the same source as `x`. Keep in mind that it might be time-consuming. With large data, when `copy = TRUE` it can speed things up to set `auto_index = TRUE`

We'll run some examples using the built-in `src` in the `dbplyr` package.

```
library(dbplyr)
```

```
##
```

```
## Attaching package: 'dbplyr'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
##      ident, sql
```

```
flights_db <- nycflights13_sqlite()
```

```
library(nycflights13)
flights2 <- flights %>%
  select(year:day, hour, tailnum, carrier)
right_join(airlines, flights2, by = "carrier")
```

```
## # A tibble: 336,776 x 7
```

```
##   carrier name          year month   day h
##   <chr>   <chr>         <int> <int> <int> <
## 1 UA      United Air Lines Inc.    2013     1     1
## 2 UA      United Air Lines Inc.    2013     1     1
## 3 AA      American Airlines Inc.    2013     1     1
## 4 B6      JetBlue Airways            2013     1     1
## 5 DL      Delta Air Lines Inc.        2013     1     1
## 6 UA      United Air Lines Inc.    2013     1     1
## 7 B6      JetBlue Airways            2013     1     1
## 8 EV      ExpressJet Airlines Inc.    2013     1     1
## 9 B6      JetBlue Airways            2013     1     1
## 10 AA     American Airlines Inc.    2013     1     1
## # ... with 336,766 more rows
```



```
flights2db <- tbl(flights_db, "flights") %>%
  select(year:day, hour, tailnum, carrier)
right_join(airlines, flights2db, by = "carrier", copy = TRUE)
```

```
## # A tibble: 336,776 x 7
```

```
##   carrier name          year month   day h
##   <chr>   <chr>         <int> <int> <int> <int>
## 1 UA      United Air Lines Inc. 2013     1     1 1
## 2 UA      United Air Lines Inc. 2013     1     1 1
## 3 AA      American Airlines Inc. 2013     1     1 1
## 4 B6      JetBlue Airways         2013     1     1 1
## 5 DL      Delta Air Lines Inc.    2013     1     1 1
## 6 UA      United Air Lines Inc. 2013     1     1 1
## 7 B6      JetBlue Airways         2013     1     1 1
## 8 EV      ExpressJet Airlines Inc. 2013     1     1 1
## 9 B6      JetBlue Airways         2013     1     1 1
## 10 AA     American Airlines Inc. 2013     1     1 1
## # ... with 336,766 more rows
```

```
flights2db_r <- collect(tbl(flights_db, "flights") %>%
  select(year:day, hour, tailnum, carrier), n = Inf)
right_join(airlines, flights2db_r, by = "carrier")
```

```
## # A tibble: 336,776 x 7
```

```
##   carrier name          year month   day h
##   <chr>   <chr>         <int> <int> <int> <int>
## 1 UA      United Air Lines Inc. 2013     1     1 1
## 2 UA      United Air Lines Inc. 2013     1     1 1
## 3 AA      American Airlines Inc. 2013     1     1 1
## 4 B6      JetBlue Airways         2013     1     1 1
## 5 DL      Delta Air Lines Inc.    2013     1     1 1
## 6 UA      United Air Lines Inc. 2013     1     1 1
## 7 B6      JetBlue Airways         2013     1     1 1
## 8 EV      ExpressJet Airlines Inc. 2013     1     1 1
## 9 B6      JetBlue Airways         2013     1     1 1
## 10 AA     American Airlines Inc. 2013     1     1 1
## # ... with 336,766 more rows
```

`dplyr` takes a lazy approach to take advantage of the database's optimized functions. It also has a few protective features:

- ▶ `nrow()` returns NA
- ▶ printing a `tbl` gives just the first 10 rows
- ▶ `tail()` gives an error

If you know SQL, there are a few other potentially helpful commands

- ▶ `compute()` and `collapse()` do not bring results back to R, and so can be used to optimize performance
- ▶ `translate_sql()` translates R expressions to SQL

# DBI

R has a nice database interface package that enables developers to, well, interface nicely with databases. As a result, there are multiple approaches (e.g., an R-Oracle interface). `dplyr`'s approach is just one; we've looked at it because it fits nicely in the framework of the course.

## In class exercise

### 0. Type

```
library(dbplyr)
example_sql <- lahman_sqlite()
```

to set up a local SQLite database of the Lahman data

1. Type `src_tbls(example_sql)` for a list of the tables in the database
2. Use a `*_join()` function to find all individuals in the Managers table who are also in the HallOfFame table. Don't bring it into R.
3. Pipe the result of the previous question to `select()` to select just the category variable.
4. Use what you've done to provide a table (in the R, not SQL, sense) of the category variable for individuals in the Hall of Fame who worked as managers